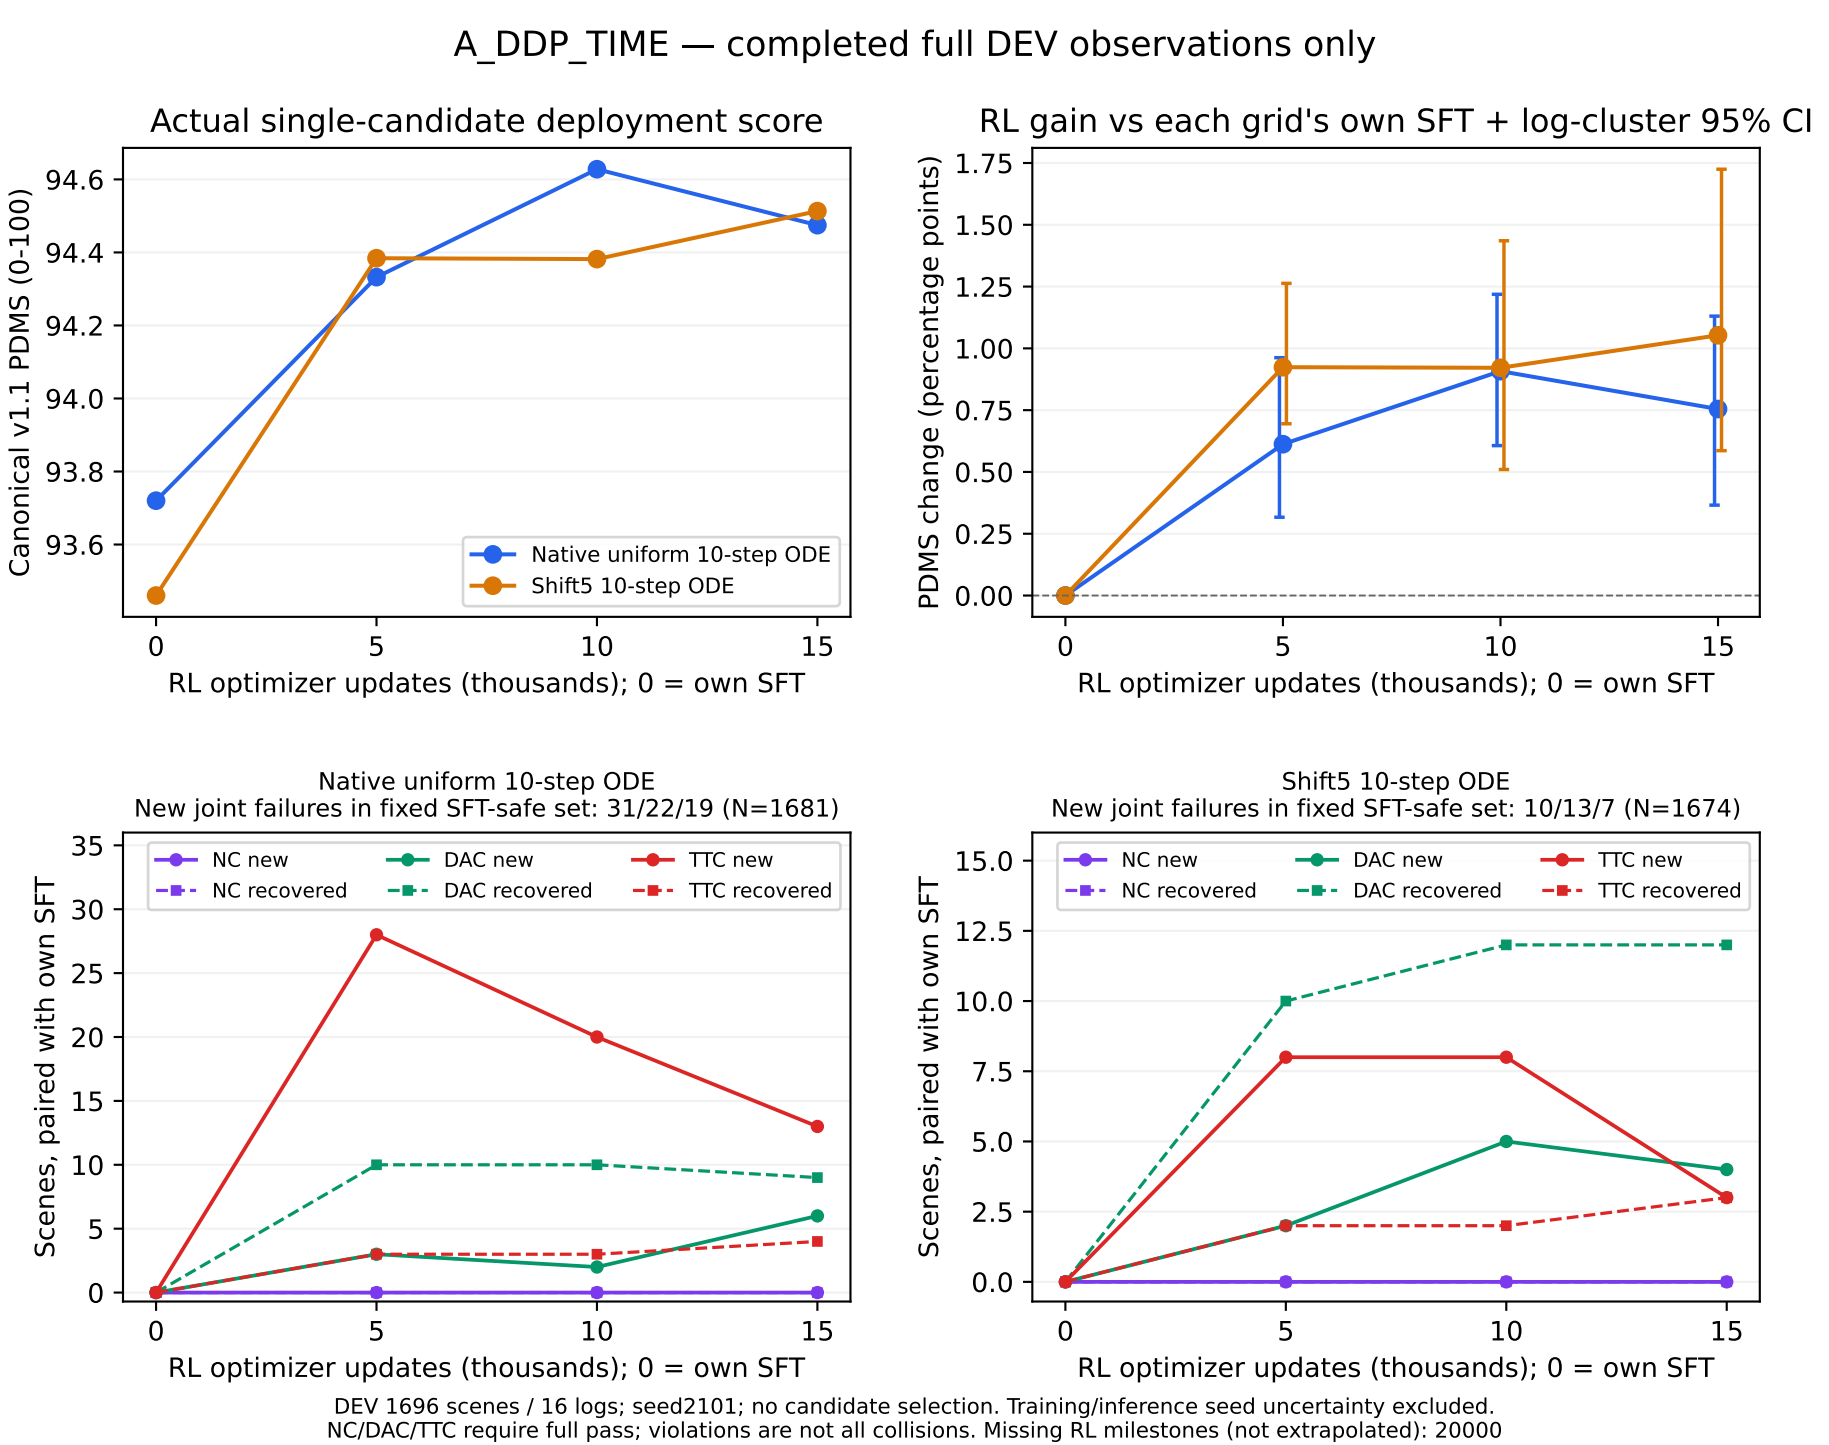
A_DDP_TIME — completed full DEV observations only
Actual single-candidate deployment score
Canonical v1.1 PDMS (0-100)
94.6
94.4
94.2
94.0
93.8
93.6
Native uniform 10-step ODE
Shift5 10-step ODE
0
5
10
15
RL optimizer updates (thousands); 0 = own SFT
RL gain vs each grid's own SFT + log-cluster 95% CI
PDMS change (percentage points)
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
RL optimizer updates (thousands); 0 = own SFT
Native uniform 10-step ODE
New joint failures in fixed SFT-safe set: 31/22/19 (N=1681)
Scenes, paired with own SFT
35
30
25
20
15
10
5
0
NC new
NC recovered
DAC new
DAC recovered
TTC new
TTC recovered
RL optimizer updates (thousands); 0 = own SFT
Shift5 10-step ODE
New joint failures in fixed SFT-safe set: 10/13/7 (N=1674)
Scenes, paired with own SFT
15.0
12.5
10.0
7.5
5.0
2.5
0.0
NC new
NC recovered
DAC new
DAC recovered
TTC new
TTC recovered
RL optimizer updates (thousands); 0 = own SFT
DEV 1696 scenes / 16 logs; seed2101; no candidate selection. Training/inference seed uncertainty excluded.
NC/DAC/TTC require full pass; violations are not all collisions. Missing RL milestones (not extrapolated): 20000