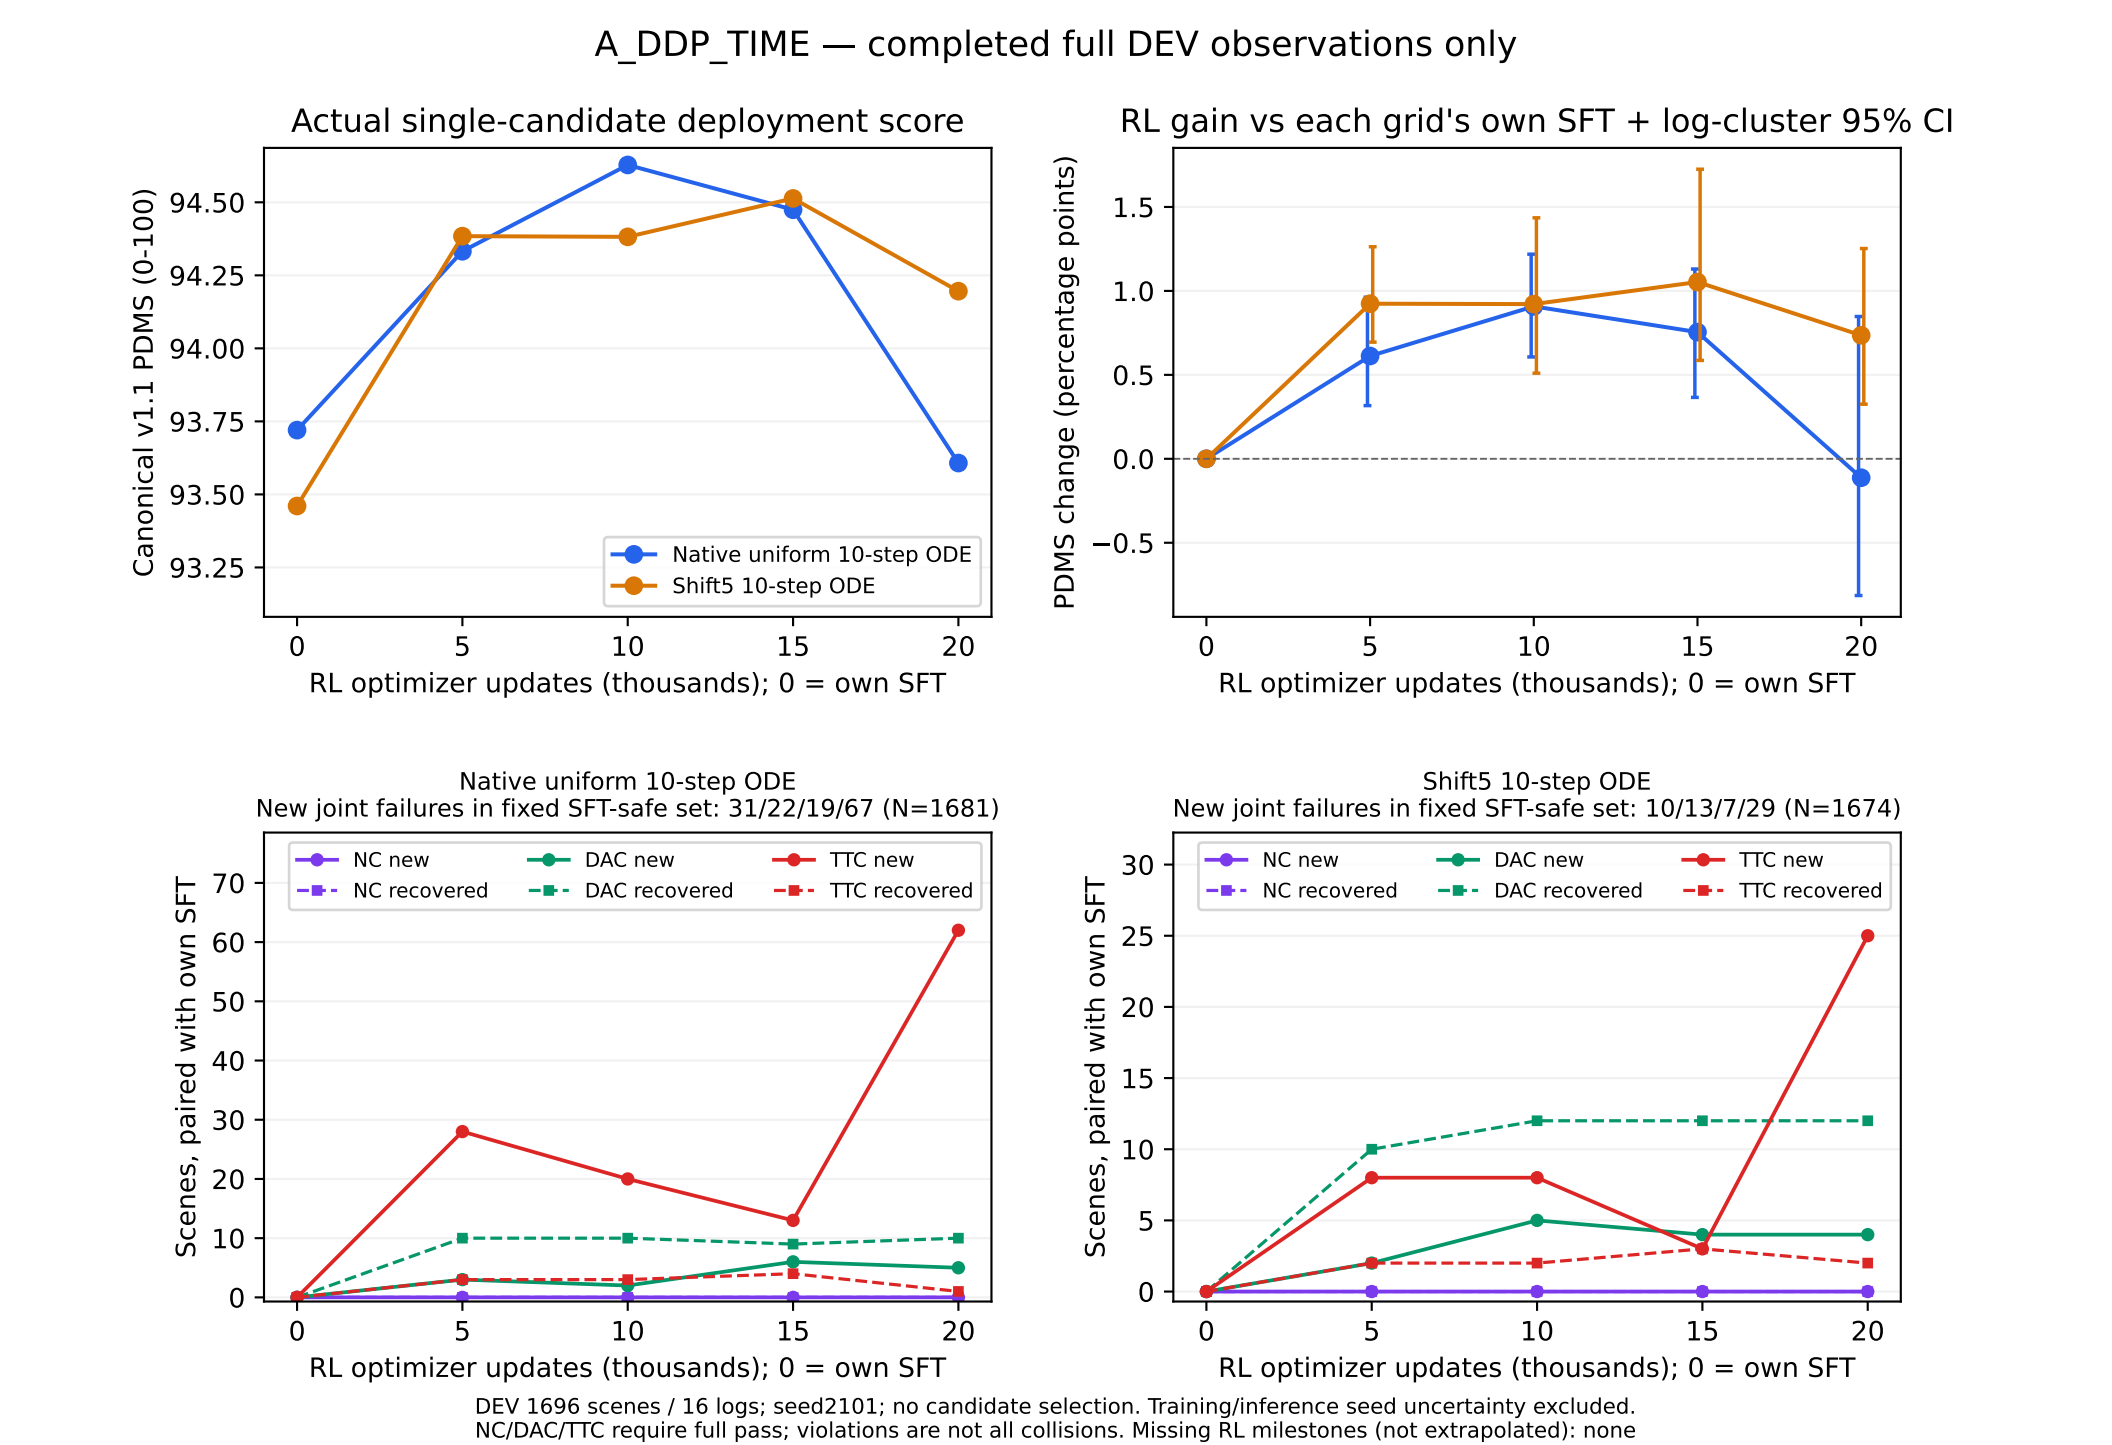
A_DDP_TIME — completed full DEV observations only
Actual single-candidate deployment score
Canonical v1.1 PDMS (0-100)
94.50
94.25
94.00
93.75
93.50
93.25
Native uniform 10-step ODE
Shift5 10-step ODE
0
5
10
15
20
RL optimizer updates (thousands); 0 = own SFT
RL gain vs each grid's own SFT + log-cluster 95% CI
PDMS change (percentage points)
1.5
1.0
0.5
0.0
−0.5
RL optimizer updates (thousands); 0 = own SFT
Native uniform 10-step ODE
New joint failures in fixed SFT-safe set: 31/22/19/67 (N=1681)
Scenes, paired with own SFT
NC new
NC recovered
DAC new
DAC recovered
TTC new
TTC recovered
RL optimizer updates (thousands); 0 = own SFT
Shift5 10-step ODE
New joint failures in fixed SFT-safe set: 10/13/7/29 (N=1674)
Scenes, paired with own SFT
RL optimizer updates (thousands); 0 = own SFT
DEV 1696 scenes / 16 logs; seed2101; no candidate selection. Training/inference seed uncertainty excluded.
NC/DAC/TTC require full pass; violations are not all collisions. Missing RL milestones (not extrapolated): none