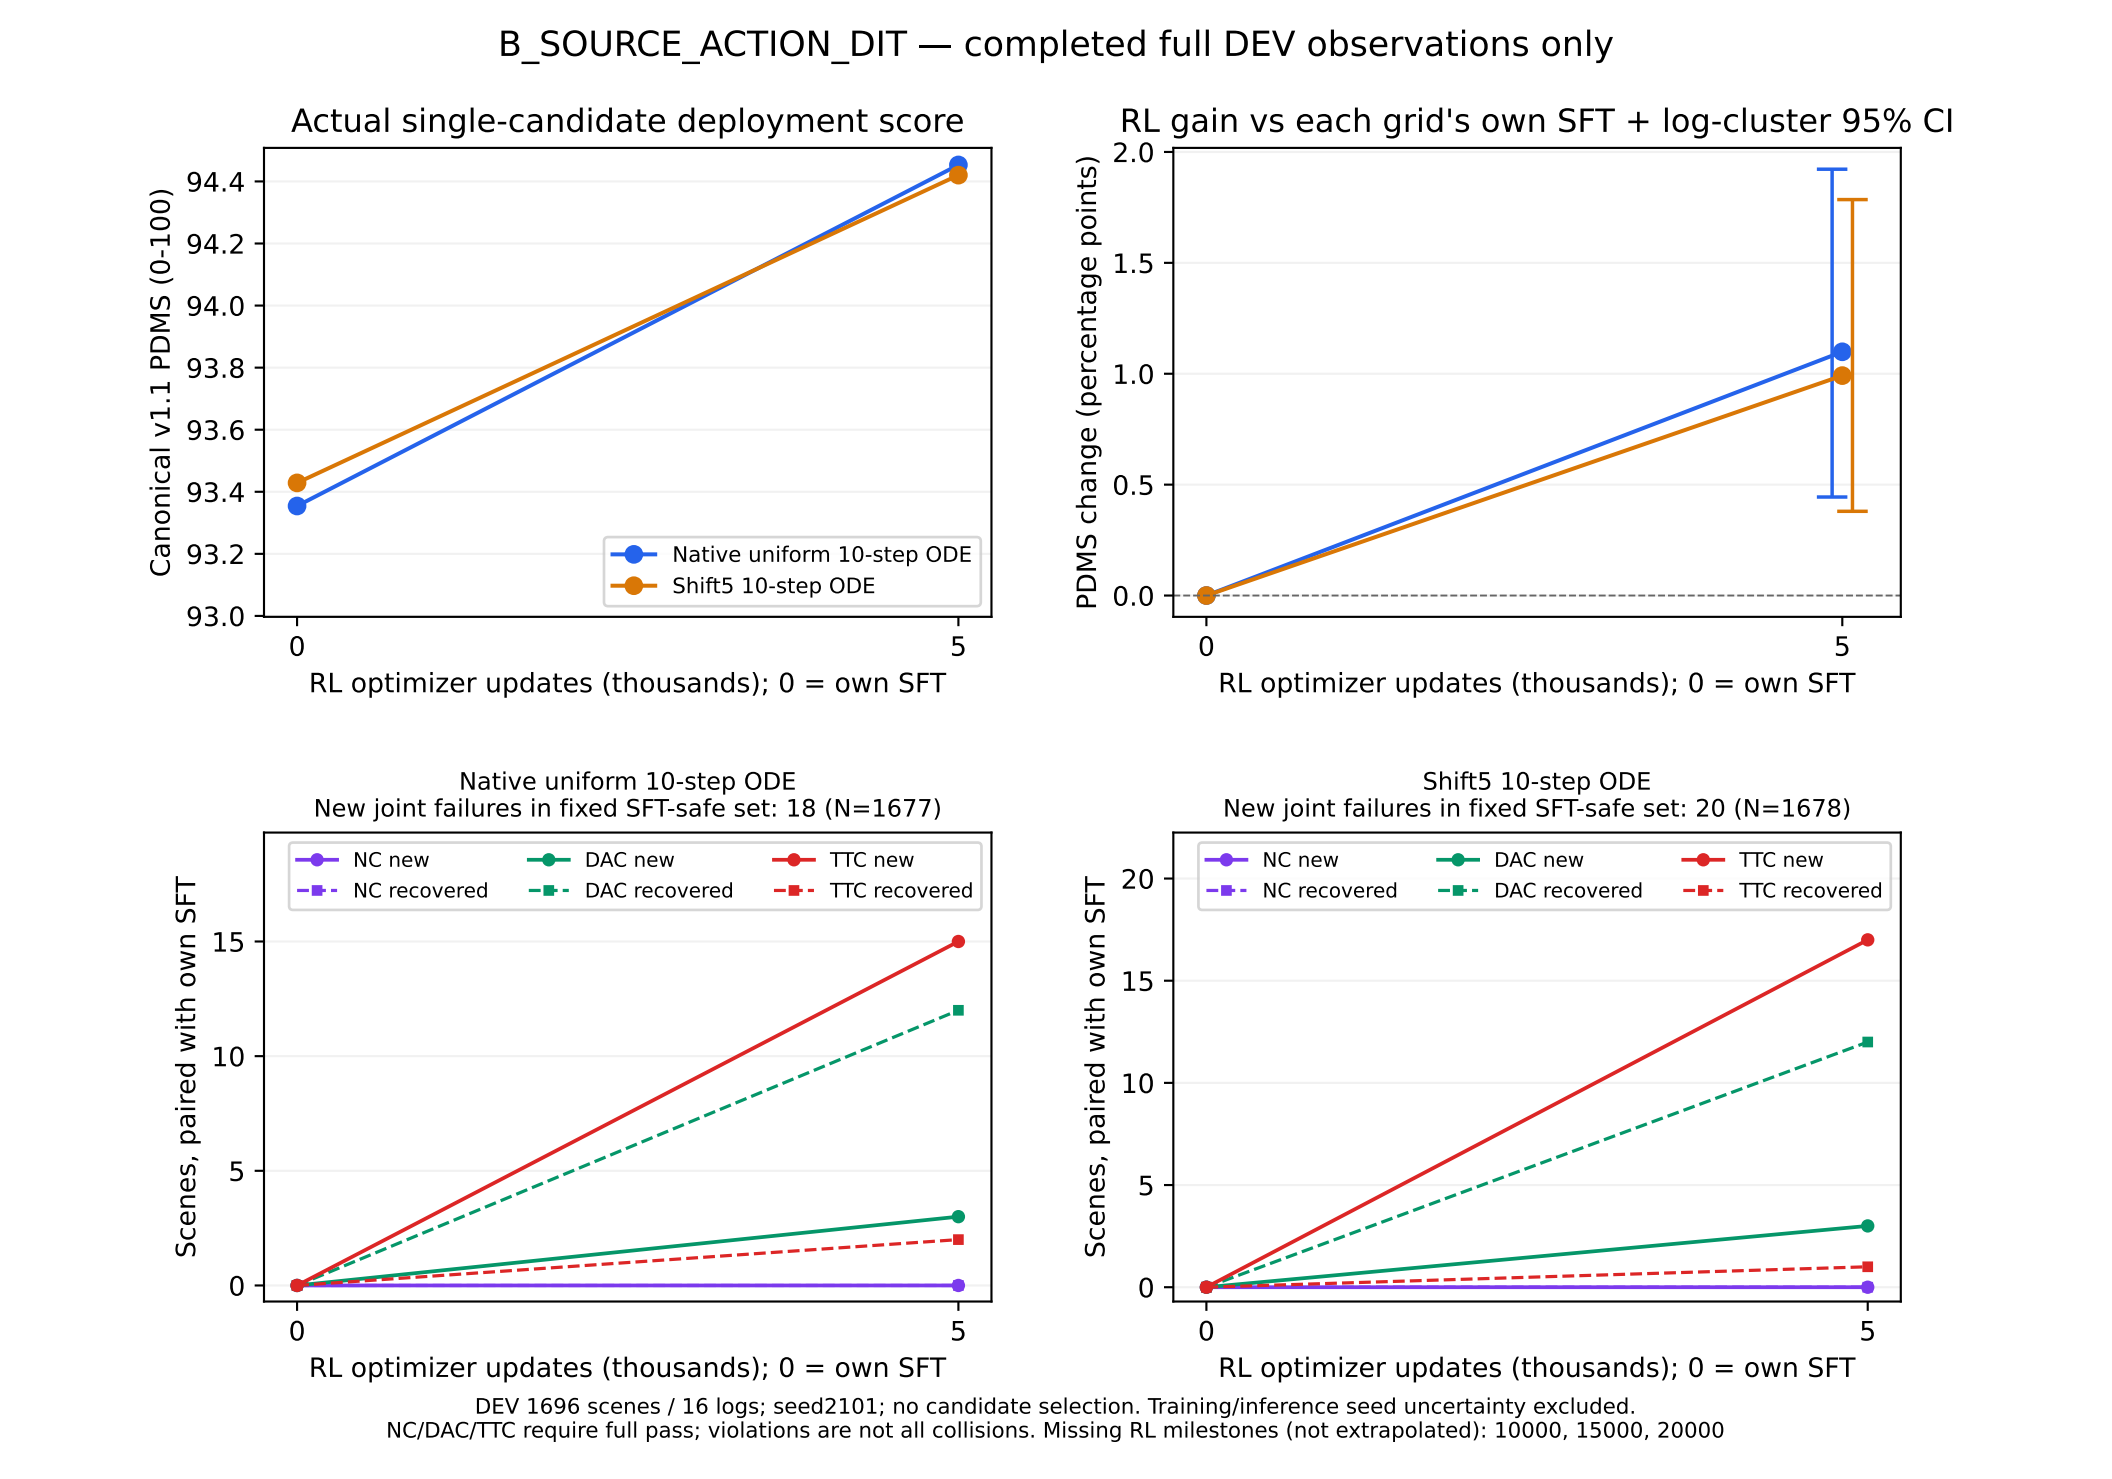
B_SOURCE_ACTION_DIT — completed full DEV observations only
Actual single-candidate deployment score
Canonical v1.1 PDMS (0-100)
94.4
94.2
94.0
93.8
93.6
93.4
93.2
93.0
Native uniform 10-step ODE
Shift5 10-step ODE
0
5
RL optimizer updates (thousands); 0 = own SFT
RL gain vs each grid's own SFT + log-cluster 95% CI
PDMS change (percentage points)
2.0
1.5
1.0
0.5
0.0
0
5
RL optimizer updates (thousands); 0 = own SFT
Native uniform 10-step ODE
New joint failures in fixed SFT-safe set: 18 (N=1677)
NC new
NC recovered
DAC new
DAC recovered
TTC new
TTC recovered
Scenes, paired with own SFT
15
10
5
0
0
5
RL optimizer updates (thousands); 0 = own SFT
Shift5 10-step ODE
New joint failures in fixed SFT-safe set: 20 (N=1678)
NC new
NC recovered
DAC new
DAC recovered
TTC new
TTC recovered
Scenes, paired with own SFT
20
15
10
5
0
0
5
RL optimizer updates (thousands); 0 = own SFT
DEV 1696 scenes / 16 logs; seed2101; no candidate selection. Training/inference seed uncertainty excluded.
NC/DAC/TTC require full pass; violations are not all collisions. Missing RL milestones (not extrapolated): 10000, 15000, 20000